Agents Can Build Benchmarks. Approvals Missed What Probing Found: 22 Cases

Pooria Arab
Independent researcher
papers@aibenchmarkindex.com

Abstract

We report 22 agent benchmarks built by AI coding agents on one shared TypeScript engine. No test job failed on the last commit of any of the 123 merged feature pull requests.

1 Introduction

Text.